

ECSP3000 Project Brief

A Reproducible Evaluation of Linux Rootkit Detectors: A Testbed and a Proof-of-Concept Rootkit

Elliott Jelbert

ejlg24@soton.ac.uk

Supervisor: Dr. BooJoong Kang

1. Project Aims

Current research into the evaluation of rootkit detectors rely on bespoke evaluation methodologies which cannot be reliably or easily reproduced by independent parties. Furthermore, as the Linux kernel is updated, the mechanisms that rootkits exploit change, as do the techniques used by detectors. This creates a need for a consistent way to evaluate rootkit detector tools continuously as new rootkit techniques are developed and new detection tools are produced.

This project aims to fulfill this need by producing a tool that provides developers and researchers with an automated and reproducible method for evaluating rootkit detectors & rootkits. This testbed will then be used alongside a modern rootkit implementation to assess the effectiveness of on-host detection techniques.

2. Project Objectives

The goal for this project is to develop a reproducible & automated testbed for linux rootkit detectors, delivering an in depth analysis of the effectiveness of on-host rootkit detectors on the modern Linux kernel.

Research Goals

- The testbed produces **reproducible** and **accurate** results in a **timely** manner.
- The testbed must allow developers to easily add new rootkit detectors or rootkits into the system
- As a stretch, design & implement a modern rootkit for the 7.x Linux kernel series.
- To combine the testbed & a modern rootkit design to produce an evaluation of the current effectiveness of on-host rootkit detection tools.

3. Limitations & Scope

This project aims to investigate on-host rootkit detection techniques only. This excludes memory forensic techniques, hypervisor-based techniques or any other out-of-band method from the scope. Similarly, only kernel-space & user-space rootkits will be used in testing. Additionally these rootkits must be public & open-source and not employ any anti-sandboxing measures which may impact results due to the virtualised environment used in testing.